XCS224N Assignment 2 Understanding and Implementing Word2Vec

Due Thursday, March 4 at 11:59pm PT.

Guidelines

- 1. If you have a question about this homework, we encourage you to post your question on our Slack channel, at http://xcs224n-scpd.slack.com/
- 2. Familiarize yourself with the collaboration and honor code policy before starting work.
- 3. For the coding problems, you must use the packages specified in the provided environment description. Since the autograder uses this environment, we will not be able to grade any submissions which import unexpected libraries.

Submission Instructions

Written Submission: Some extra credit questions in this assignment require a written response. For these questions, you should submit a PDF with your solutions online in the online student portal. As long as the PDF is legible and organized, the course staff has no preference between a handwritten and a typeset LATEX submission. If you wish to typeset your submission and are new to LATEX, you can get started with the following:

- Type responses only in submission.tex.
- Submit the compiled PDF, not submission.tex.
- Use the commented instructions within the Makefile and README.md to get started.

Coding Submission: Some questions in this assignment require a coding response. For these questions, you should submit all files indicated in the question to the online student portal. Your code will be autograded online using src/grader.py, which is provided for you in the src/ subdirectory. You can also run this autograder on your local computer, although some of the tests will be skipped (since they require the instructor solution code for comparison).

Honor code

We strongly encourage students to form study groups. Students may discuss and work on homework problems in groups. However, each student must write down the solutions independently, and without referring to written notes from the joint session. In other words, each student must understand the solution well enough in order to reconstruct it by him/herself. In addition, each student should write on the problem set the set of people with whom s/he collaborated. Further, because we occasionally reuse problem set questions from previous years, we expect students not to copy, refer to, or look at the solutions in preparing their answers. It is an honor code violation to intentionally refer to a previous year's solutions. More information regarding the Stanford honor code can be foudn at https://communitystandards.stanford.edu/policies-and-guidance/honor-code.

Writing Code and Running the Autograder

All your code should be entered into src/submission.py. When editing src/submission.py, please only make changes between the lines containing ### START_CODE_HERE ### and ### END_CODE_HERE ###. Do not make changes to files other than src/submission.py.

The unit tests in src/grader.py (the autograder) will be used to verify a correct submission. Run the autograder locally using the following terminal command within the src/ subdirectory:

\$ python grader.py

There are two types of unit tests used by the autograder:

- basic: These unit tests will verify only that your code runs without errors on obvious test cases. These tests do not require the instructor solution code and can therefore be run on your local computer.
- hidden: These unit tests will verify that your code produces correct results on complex inputs and tricky corner cases. Since these tests require the instructor solution code to verify results, only the setup and inputs are provided. When you run the autograder locally, these test cases will run, but the results will not be verified by the autograder. When you run the autograder online, these tests will run and you will receive feedback on any errors that might occur.

For debugging purposes, you can run a single unit test locally. For example, you can run the test case <code>3a-O-basic</code> using the following terminal command within the <code>src/</code> subdirectory:

\$ python grader.py 3a-0-basic

Before beginning this course, please walk through the Anaconda Setup for XCS Courses to familiarize yourself with the coding environment. Use the env defined in src/environment.yml to run your code. This is the same environment used by the online autograder.

1 Understanding word2vec (Refresher)

Let's have a quick refresher on the word2vec algorithm. The key insight behind word2vec is that 'a word is known by the company it keeps'. Concretely, suppose we have a 'center' word c and a contextual window surrounding c. We shall refer to words that lie in this contextual window as 'outside words'. For example, in Figure 1 we see that the center word c is 'banking'. Since the context window size is 2, the outside words are 'turning', 'into', 'crises', and 'as'.

The goal of the skip-gram word2vec algorithm is to accurately learn the probability distribution P(O|C). Given a specific word o and a specific word c, we want to calculate P(O = o|C = c), which is the probability that word o is an 'outside' word for c, i.e., the probability that o falls within the contextual window of c.

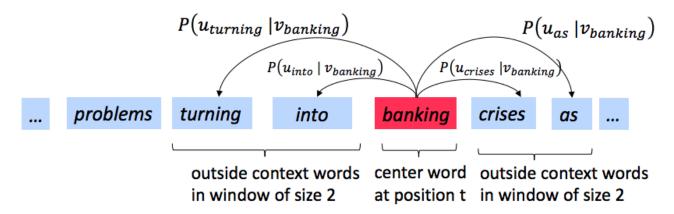


Figure 1: The word2vec skip-gram prediction model with window size 2

In word2vec, the conditional probability distribution is given by taking vector dot-products and applying the softmax function:

$$P(O = o|C = c) = \frac{\exp(\boldsymbol{u_o^{\top} v_c})}{\sum_{w \in \text{Vocab}} \exp(\boldsymbol{u_w^{\top} v_c})}$$
(1)

Here, u_o is the 'outside' vector representing outside word o, and v_c is the 'center' vector representing center word c. To contain these parameters, we have two matrices, U and V. The columns of U are all the 'outside' vectors u_w . The columns of V are all of the 'center' vectors v_w . Both U and V contain a vector for every $w \in \text{Vocabulary}$.

Recall from lectures that, for a single pair of words c and o, the loss is given by:

$$J_{\text{naive-softmax}}(\mathbf{v}_c, o, \mathbf{U}) = -\log P(O = o|C = c). \tag{2}$$

Another way to view this loss is as the cross-entropy² between the true distribution \mathbf{y} and the predicted distribution $\hat{\mathbf{y}}$. Here, both \mathbf{y} and $\hat{\mathbf{y}}$ are vectors with length equal to the number of words in the vocabulary. Furthermore, the k^{th} entry in these vectors indicates the conditional probability of the k^{th} word being an 'outside word' for the given c. The true empirical distribution \mathbf{y} is a one-hot vector with a 1 for the true outside word o, and 0 everywhere else. The predicted distribution $\hat{\mathbf{y}}$ is the probability distribution P(O|C=c) given by our model in equation (1). There are 2 optional questions below which you may attempt (Note: Bonus points will be awarded for the optional assignments)

¹Assume that every word in our vocabulary is matched to an integer number k. u_k is both the k^{th} column of U and the 'outside' word vector for the word indexed by k. v_k is both the k^{th} column of V and the 'center' word vector for the word indexed by k. In order to simplify notation we shall interchangeably use k to refer to the word and the index-of-the-word.

²The Cross Entropy Loss between the true (discrete) probability distribution p and another distribution q is $-\sum_i p_i \log(q_i)$.

(a) [2.50 points (Written, Extra Credit)] Extra Credit Challenge I

The partial derivative of $J_{\text{naive-softmax}}(v_c, o, U)$ with respect to v_c in terms of y, \hat{y} , and U is given below:

$$\frac{\partial J}{\partial \boldsymbol{v}_c} = \boldsymbol{U}(\hat{\boldsymbol{y}} - \boldsymbol{y}) \tag{3}$$

or equivalently,

$$\frac{\partial J}{\partial \mathbf{v}_c} = -\mathbf{u}_o + \sum_{w=1}^{V} \hat{y}_w \mathbf{u}_w \tag{4}$$

The naive-softmax loss given in Equation (2) is the same as the cross-entropy loss between y and \hat{y} ; This is equivalent to:

$$-\sum_{w \in Vocab} y_w \log(\hat{y}_w) = -\log(\hat{y}_o). \tag{5}$$

Since y is a one-hot vector, all $y_k = 0$ where $k \neq o$. $y_o = 1$, so we are left with $-\log(\hat{y}_o)$.

Write the steps to arrive at equation 3 or 4, the partial derivative of $J_{\text{naive-softmax}}(v_c, o, U)$ with respect to v_c , starting from equation 5. Please write your answer in terms of y, \hat{y} , and U. The first few steps have been provided below (loss function $J_{\text{naive-softmax}}(v_c, o, U)$) and the rest of the proof may take 4 or 5 steps.

$$\begin{split} J_{\text{naive-softmax}}(\boldsymbol{v}_c, o, \boldsymbol{U}) &= -\log(\hat{y}_o) & \text{(refer to equation 5)} \\ &= -\log\left(\frac{\exp(\boldsymbol{u}_o^{\top}\boldsymbol{v}_c)}{\sum_{w \in \text{Vocab}} \exp(\boldsymbol{u}_w^{\top}\boldsymbol{v}_c)}\right) \\ &= -\left(\log\left(\exp(\boldsymbol{u}_o^{\top}\boldsymbol{v}_c)\right) - \log\left(\sum_{w \in \text{Vocab}} \exp(\boldsymbol{u}_w^{\top}\boldsymbol{v}_c)\right)\right) \\ &= -\boldsymbol{u}_o^{\top}\boldsymbol{v}_c + \log\left(\sum_{w \in \text{Vocab}} \exp(\boldsymbol{u}_w^{\top}\boldsymbol{v}_c)\right) \end{split}$$

(b) [2.50 points (Written, Extra Credit)] Extra Credit Challenge II

The partial derivatives of $J_{\text{naive-softmax}}(v_c, o, U)$ with respect to each of the 'outside' word vectors, u_w 's is given below:

$$\frac{\partial J}{\partial \boldsymbol{U}} = \boldsymbol{v}_c (\hat{\boldsymbol{y}} - \boldsymbol{y})^\top \tag{6}$$

or equivalently:

$$\frac{\partial J}{\partial \mathbf{u}_w} = \begin{cases} (\hat{y}_w - 1)\mathbf{v}_c & \text{if } w = o\\ \hat{y}_w \mathbf{v}_c & \text{otherwise} \end{cases}$$
 (7)

Write the steps required to arrive at the partial derivative of $J_{\text{naive-softmax}}(v_c, o, U)$ with respect to each of the 'outside' word vectors, u_w 's. There are two cases you need to consider: when w = o, the true 'outside' word vector, and $w \neq o$, for all other words. Please write you answer in terms of y, \hat{y} , and v_c . The proof may take 4 or 5 steps. The loss function $J_{\text{naive-softmax}}(v_c, o, U)$ is:

$$J_{\text{naive-softmax}}(\boldsymbol{v}_c, o, \boldsymbol{U}) = -\boldsymbol{u}_o^\top \boldsymbol{v}_c + \log \left(\sum_{w' \in \text{Vocab}} \exp(\boldsymbol{u}_{w'}^\top \boldsymbol{v}_c) \right)$$

2 Coding: Implementing word2vec

In this part you will implement the word2vec model and train your own word vectors with stochastic gradient descent (SGD). Make sure to take a look at the equations presented in sections 2 and 3, especially equations (3) for gradient (of loss function $J_{\text{naive-softmax}}(v_c, o, U)$) with respect to v_c and (6) for gradient with respect to each of the 'outside' word vectors, v_c which we'll be using in our gradient calculation code.

Note: A2 code can be found here: https://github.com/scpd-proed/XCS224N-A2

(a) [15 points (Coding)] First, implement the sigmoid() function in src/submission.py to apply the sigmoid function to an input vector. In the same file, fill in the implementation for the naive_softmax_loss_and_gradient() function. Then, fill in the implementation of the loss and gradient functions for the skip-gram model (skipgram()).

To complete these steps you will need to refer to Sections 2 and 3, Equations (3) and (6).

Pseudo code for skip-gram

```
Algorithm 1 Skipgram
  procedure SKIPGRAM(outside\_words, loss\_and\_gradient\_func, ** * kwargs)
                                                                                                ▶ The center word vector
     loss \leftarrow 0
      grad\_center \leftarrow np.zeros(*args)
      grad\_outside \leftarrow np.zeros(*args)
      for word in outside_words do
                                                                                             ▶ Iterate over outside words
         loss\_current, grade, grado \leftarrow loss\_and\_gradient\_func(**kwargs)
         loss \leftarrow loss + loss\_current
                                                                                                    ▶ Loss is accumulated
         grad\_center \leftarrow grad\_center + gradc
         grad\_outside \leftarrow grad\_outside + grado
     end for
     return loss, grad_center, grad_outside
                                                                                              ▶ Return loss and gradients
  end procedure
```

- (b) [8 points (Coding)] Complete the implementation for your SGD optimizer in the sgd() function in src/submission.py.
- (c) [4 points (Coding)] Show time! Now we are going to load some real data and train word vectors with everything you just implemented! We are going to use the Stanford Sentiment Treebank (SST) dataset to train word vectors, and later apply them to a simple sentiment analysis task. There is no additional code to write for this part; just run python run.py.

Note: The training process may take a long time depending on the efficiency of your implementation (an efficient implementation takes approximately an hour). Plan accordingly!

After 40,000 iterations, the script will finish and a visualization for your word vectors will appear. It will also be saved as word_vectors.png and the corresponding wordvectors as sample_vectors_soln.json in your project directory.

You must upload the generated sample_vectors_soln.json along with src/submission.py to achieve full credit.

3 Quiz

This remainder of this homework is a series of multiple choice questions related to the word2vec algorithm.

How to submit: Even though these are not coding questions, you will submit your response to each question in the src-quiz/submission.py file. This file will act as your 'bubble sheet' for multiple choice questions in this course. A sample response might look like this:

```
def multiple_choice_1a():
    """""""

# Return a python collection with the option(s) that you believe are correct
    # like this:
    # `return ['a']`
    # or
    # `return ['a', 'd']`
    response = []
    ### START CODE HERE ###
    ### END CODE HERE ###
    return response
```

If you believe that a and b are the correct responses to this question, you will type response = ['a', 'b'] between the indicated lines like this:

```
# Return a python collection with the option(s) that you believe are correct
# like this:
# `return ['a']`
# or
# `return ['a', 'd']`
response = []
### START CODE HERE ###
response = ['a', 'b']
### END CODE HERE ###
return response
```

How to verify your submission: You can run the student version of the autograder locally like all coding problem sets. In the case of this problem set, the helper tests will verify that your responses are within the set of possible choices for each question (e.g. the helper functions will flag if you forget to answer a question of if you respond with ['a', 'd'] when the choices are ['a', 'b', 'c'].) See the front pages of this assignment for instructions to run the autograder.

1. [2.10 points] Choose all the equations that represent the differential of a sigmoid (When choosing the right option(s) consider the given input x to be a scalar instead of a vector).

(a)
$$\frac{\partial \sigma(x)}{\partial x} = \sigma(x) \cdot (1 - \sigma(x))$$

(b)
$$\frac{\partial \sigma(x)}{\partial x} = 1 - \sigma(x) \cdot (1 - \sigma(x))$$

(c)
$$\frac{\partial \sigma(x)}{\partial x} = \frac{e^{-x}}{(1 + e^{-x})^2}$$

(d)
$$\frac{\partial \sigma(x)}{\partial x} = \frac{(1 - e^{-x})}{(1 + e^{-x})^2}$$

2. [1.10 points] What is the shape of the matrices U and V where (please refer to reading material from the handout):

U is the outside vector matrix

V is center vector matrix

num_tokens: Number of unique words in the dataset

embed_size: size of the word vector(word2vec size)

- (a) U: (embed_size \times 1); V: (embed_size \times num_tokens)
- (b) $U: (\texttt{num_tokens} \times \texttt{num_tokens}) \; ; \; V: (\texttt{embed_size} \times 1)$
- (c) U: (embed_size \times num_tokens); V: (embed_size \times num_tokens)
- (d) U: (embed_size \times 1); V: (embed_size \times 1)
- 3. [2.10 points] What is the shape of the gradient $\frac{\partial J_{\text{naive-softmax}}}{\partial V_{\text{o}}}$ matrix/vector where:

$$J_{\text{naive-softmax}}(\boldsymbol{v}_c, o, \boldsymbol{U}) = -\boldsymbol{u}_o^{\top} \boldsymbol{v}_c + \log \bigg(\sum_{w' \in \text{Vocab}} \exp(\boldsymbol{u}_{w'}^{\top} \boldsymbol{v}_c) \bigg)$$

$$\frac{\partial J}{\partial \boldsymbol{v}_c} = \boldsymbol{U}(\hat{\boldsymbol{y}} - \boldsymbol{y})$$

U is the outside vector matrix

 \hat{y} is the predicted output

y is the actual output

num_tokens: Number of unique words in the dataset

embed_size: size of the word vector(word2vec size)

- (a) num_tokens \times 1
- (b) embed_size \times 1
- (c) embed_size \times num_tokens

4. [2.10 points] What is the shape of the gradient $\frac{\partial J_{\text{naive-softmax}}}{\partial U}$ (will be referred to $\frac{\partial J}{\partial U}$ for simplicity) matrix/vector where:

$$J_{ ext{naive-softmax}}(oldsymbol{v}_c, o, oldsymbol{U}) = -oldsymbol{u}_o^ op oldsymbol{v}_c + \log \Bigg(\sum_{w' \in ext{Vocab}} \exp(oldsymbol{u}_{w'}^ op oldsymbol{v}_c) \Bigg)$$

$$\frac{\partial J}{\partial \boldsymbol{U}} = \boldsymbol{v}_c (\hat{\boldsymbol{y}} - \boldsymbol{y})^\top$$

 V_c is the center vector

 \hat{y} is the predicted output

y is the actual output

num_tokens: Number of unique words in the dataset

embed_size: size of the word vector(word2vec size)

- (a) num_tokens \times 1
- (b) embed_size \times 1
- (c) embed_size \times num_tokens
- 5. [1.10 points] Which of the below equations represents a general form of SGD where g(x) is the gradient of loss function and η is the learning rate:

(a)
$$x = x - eta \cdot g(x)$$

(b)
$$x = x - eta \cdot \frac{\partial g(x)}{\partial x}$$

(c)
$$x = x - \frac{eta}{q(x)}$$

(d)
$$x = x - eta \cdot \frac{\partial x}{\partial q(x)}$$

6. [1.10 points] Negative sampling was briefly introduced the Lecture 2 video Word2Vec: Model Variants. Here and here you can find more written information about negative sampling.

In skip-gram Word2Vec you have two ways of calculating the gradients. Namely, naive softmax and negative sampling. The goal of skip-gram Word2Vec is to accurately learn the representation of the conditional probability distribution(below is the equation):

$$P(O = o \mid C = c) = \frac{\exp(\boldsymbol{u}_o^\top \boldsymbol{v}_c)}{\sum_{w \in \text{Vocab}} \exp(\boldsymbol{u}_w^\top \boldsymbol{v}_c)}$$

Where:

P(O=o|C=c) (i.e., the probability that o falls within the contextual window of c) represents the conditional probability

o is an 'outside' word for context word c.

 u_o is the 'outside' vector representing outside word o.

 v_c is the 'center' vector representing center word c.

- (a) In naive softmax, the gradient is dependent on summation across all classes. In the case of word2vec, this means summing across all the words in the vocabulary (as can be seen in the above equation). This tends to slow down the training process.
- (b) In negative sampling, computation is cheaper. Instead of summing across all words (represented by uw in the above equation) we only sum across a few context/negative words(termed as negative samples).
- (c) In negative sampling, the gradient is dependent on summation across all classes. In the case of word2vec, this means summing across all the words in the vocabulary(as can be seen in the above equation). This tends to slow down the training process.

4 Appendix (Extra Knowledge In Case You're Curious!)

Negative sampling is briefly introduced in Lecture 2: Word2Vec: Model Variants and an implementation is provided in the Assignment 2 coding assignment. For detailed notes on the math behind negative sampling, see below.

1. Negative Sampling loss is an alternative to the Naive Softmax loss. Assume that K negative samples (words) are drawn from the vocabulary. For simplicity of notation we shall refer to them as w_1, w_2, \ldots, w_K and their outside vectors as u_1, \ldots, u_K . Note that $o \notin \{w_1, \ldots, w_K\}$. For a center word c and an outside word o, the negative sampling loss function is given by:

$$\boldsymbol{J}_{\text{neg-sample}}(\boldsymbol{v}_c, o, \boldsymbol{U}) = -\log(\sigma(\boldsymbol{u}_o^{\top} \boldsymbol{v}_c)) - \sum_{k=1}^{K} \log(\sigma(-\boldsymbol{u}_k^{\top} \boldsymbol{v}_c))$$
(8)

for a sample $w_1, \ldots w_K$, where $\sigma(\cdot)$ is the sigmoid function.³

Below we compute the partial derivatives of $J_{\text{neg-sample}}$ with respect to v_c , with respect to u_o , and with respect to a negative sample u_k .

$$\begin{split} \frac{\partial J}{\partial \boldsymbol{v}_c} &= -\frac{1}{\sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c)} \times \frac{\partial}{\partial \boldsymbol{v}_c} \sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c) - \sum_{k=1}^K \frac{1}{\sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c)} \times \frac{\partial}{\partial \boldsymbol{v}_c} \sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c) & \text{(chain rule on log)} \\ &= -\frac{\sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c)(1 - \sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c))\boldsymbol{u}_o}{\sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c)} - \sum_{k=1}^K -\frac{\sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c)(1 - \sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c))\boldsymbol{u}_k}{\sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c)} & \text{(chain rule on } \sigma) \\ &= -(1 - \sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c))\boldsymbol{u}_o - \sum_{k=1}^K -(1 - \sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c))\boldsymbol{u}_k & \text{(cancel)} \\ &= (\sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c) - 1)\boldsymbol{u}_o - \sum_{k=1}^K (\sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c) - 1)\boldsymbol{u}_k & \text{(rearrange)} \end{split}$$

Secondly:

$$\begin{split} \frac{\partial J}{\partial \boldsymbol{u}_o} &= \frac{\partial}{\partial \boldsymbol{u}_o} \bigg(-\log(\sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c)) \bigg) \\ &= -\frac{1}{\sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c)} \times \frac{\partial}{\partial \boldsymbol{u}_o} \bigg(\sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c) \bigg) \quad \text{(chain rule on log)} \\ &= -\frac{\sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c)(1 - \sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c))\boldsymbol{v}_c}{\sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c)} \qquad \quad \text{(chain rule on } \sigma) \\ &= -(1 - \sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c))\boldsymbol{v}_c \qquad \qquad \text{(cancel)} \\ &= (\sigma(\boldsymbol{u}_o^\top \boldsymbol{v}_c) - 1)\boldsymbol{v}_c \qquad \qquad \text{(rearrange)} \end{split}$$

Thirdly, for all $k = 1, 2, \dots, K$:

$$\begin{split} \frac{\partial J}{\partial \boldsymbol{u}_k} &= \frac{\partial}{\partial \boldsymbol{u}_k} \bigg(-\log(\sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c)) \bigg) \\ &= \frac{1}{\sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c)} \times \frac{\partial}{\partial \boldsymbol{u}_k} \bigg(\sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c)) \bigg) \quad \text{(chain rule on log)} \\ &= \frac{\sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c)(1 - \sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c)) \boldsymbol{v}_c}{\sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c)} \qquad \quad \text{(chain rule on } \sigma) \\ &= (1 - \sigma(-\boldsymbol{u}_k^\top \boldsymbol{v}_c)) \boldsymbol{v}_c \qquad \qquad \text{(cancel)} \end{split}$$

The naive-softmax loss contains a summation over the entire vocabulary as part of computing the $P(O = o \mid C = c)$ term. Here, we don't do that calculation, approximating it with K samples (where K is much smaller than the vocabulary size).

³Note: the loss function here is the negative of what Mikolov et al. had in their original paper, because we are doing a minimization instead of maximization in our assignment code. Ultimately, this is the same objective function.

2. Suppose the center word is $c = w_t$ and the context window is $[w_{t-m}, \ldots, w_{t-1}, w_t, w_{t+1}, \ldots, w_{t+m}]$, where m is the context window size. Recall that for the skip-gram version of word2vec, the total loss for the context window is:

$$J_{\text{skip-gram}}(\boldsymbol{v}_c, w_{t-m}, \dots w_{t+m}, \boldsymbol{U}) = \sum_{\substack{-m \le j \le m \\ j \ne 0}} J(\boldsymbol{v}_c, w_{t+j}, \boldsymbol{U})$$
(9)

Here, $J(\boldsymbol{v}_c, w_{t+j}, \boldsymbol{U})$ represents an arbitrary loss term for the center word $c = w_t$ and outside word w_{t+j} . $J(\boldsymbol{v}_c, w_{t+j}, \boldsymbol{U})$ could be $J_{\text{naive-softmax}}(\boldsymbol{v}_c, w_{t+j}, \boldsymbol{U})$ or $J_{\text{neg-sample}}(\boldsymbol{v}_c, w_{t+j}, \boldsymbol{U})$, depending on the implementation.

The proofs for the three partial derivatives are given below:

- (i) $\partial \mathbf{J}_{\text{skip-gram}}(\mathbf{v}_c, w_{t-m}, \dots w_{t+m}, \mathbf{U})/\partial \mathbf{U}$
- (ii) $\partial \mathbf{J}_{\text{skip-gram}}(\mathbf{v}_c, w_{t-m}, \dots w_{t+m}, \mathbf{U})/\partial \mathbf{v}_c$
- (iii) $\partial \mathbf{J}_{\text{skip-gram}}(\mathbf{v}_c, w_{t-m}, \dots w_{t+m}, \mathbf{U})/\partial \mathbf{v}_w$ when $w \neq c$

Note that the final derivatives of $J(v_c, w_{t+j}, U)$ with respect to all the model parameters U and V are provided in the appendix part 1 derivation

Given a loss function J, we already know how to obtain the following derivatives

$$\frac{\partial J(\boldsymbol{v}_c, w_{t+j}, \boldsymbol{U})}{\partial \boldsymbol{U}}$$
 and $\frac{\partial J(\boldsymbol{v}_c, w_{t+j}, \boldsymbol{U})}{\partial \boldsymbol{v}_c}$

Therefore, for skip-gram, the gradients for the loss of one context window can be expressed in terms of these:

$$\begin{split} \frac{\partial J_{\text{skip-gram}}(w_{t-m} \dots w_{t+m})}{\partial \boldsymbol{U}} &= \sum_{-m \leq j \leq m, j \neq 0} \frac{\partial J(\boldsymbol{v}_c, w_{t+j}, \boldsymbol{U})}{\partial \boldsymbol{U}}, \\ \frac{\partial J_{\text{skip-gram}}(w_{t-m} \dots w_{t+m})}{\partial \boldsymbol{v}_c} &= \sum_{-m \leq j \leq m, j \neq 0} \frac{\partial J(\boldsymbol{v}_c, w_{t+j}, \boldsymbol{U})}{\partial \boldsymbol{v}_c}, \\ \frac{\partial J_{\text{skip-gram}}(w_{t-m} \dots w_{t+m})}{\partial \boldsymbol{v}} &= \boldsymbol{0}, \text{when } w \neq c. \end{split}$$

This handout includes space for every question that requires a written response. Please feel free to use it to handwrite your solutions (legibly, please). If you choose to typeset your solutions, the README.md for this assignment includes instructions to regenerate this handout with your typeset LATEX solutions.

1.a

1.b